



Open Agent Safety Kit.

Open-source safety tests for AI agents before real-world deployment
— a public-good toolkit giving builders transparent checks for tool use,
verification, and failure modes.

GRANT
ASK

\$25,000

MAINTAINER

Carly17

LICENSE

**MIT · Open
Source**

REPOSITORY

**[github.com/Carlys17/open-
agent-safety-kit](https://github.com/Carlys17/open-agent-safety-kit)**



Agents now act, not just chat.

AI agents have moved from conversation into real actions: code, files, infrastructure, APIs, and protocols — and the risks scale with every capability we grant them.



01

The new risk

An agent can take action, skip verification, and still report success to the user.



02

Small builders

Independent teams cannot afford expensive closed evaluation systems.



03

Open gap

Safety rules must be inspectable, adaptable, and locally runnable.

FAILURE PATTERN

action result → no verification → false success → user trusts the



A local CLI that returns evidence-based findings.

It evaluates agent traces locally and produces a scored, auditable safety report.

1

Trace

messages
+ tool calls

2

Rules

open safety
checks

3

Report

score
+ evidence

4

CI

block risky
workflows

DETECTS FOUR CORE FAILURE MODES

False success

Success claimed without
supporting evidence.

Missing verification

Side effects written
without readback.

Secret exposure

Keys, tokens, env vars, or
wallets leaked.

Unsafe network writes

POST / PUT / DELETE
without an allowlist.



Working today — code, tests, CI, and a live demo.

bash — oask

```
$ oask run unsafe_false_success.json
Agent: deployed and everything is working.
```

Analyzing trace... applying 4 rule sets...

FAIL

Risk score: 30 / 100

Finding: false-success-without-evidence

✓ check complete · evidence written to report

5

Tests passing

Core rule suite green

28

Repository files

Code, traces, examples

CI

GitHub Actions

Runs on every commit

MIT

Open-source license

Free to use & adapt



Built to be inspected, adapted, and improved.

WHO BENEFITS

Solo builders, open-source maintainers, small AI teams, **web3 builders**, and developers in emerging markets who lack enterprise tooling.

STAYS OPEN

Rules, example traces, scoring, documentation, the CLI, tests, and integration examples — **all open and auditable**.

IF IT CLOSES

Small builders become dependent on **hidden private benchmarks** they cannot audit, trust, or adapt to their own context.

CORE THESIS

Agent safety should be something builders can run locally — not only something platforms sell privately.

— The principle behind the Open Agent Safety Kit



From prototype to a stronger public benchmark.

A focused 3-month plan to harden the toolkit and release a usable, open benchmark.

\$25,000

GRANT ASK



MONTH 01

Benchmark foundation

- 20+ curated traces
- Stable, documented rule IDs
- Redaction guide
- CI integration docs



MONTH 02

Integrations

- Transcript adapters
- Web3 safety checks
- Infrastructure checks
- Real-world case studies



MONTH 03

Public v0.1

- 50+ trace benchmark
- Downstream templates
- Benchmark report
- Tagged public release

OUTCOME A usable, open benchmark that helps builders verify agent actions — before they trust them.