


Open Agent Safety Kit

Open-source safety tests for AI agents before real-world deployment.

Public-good toolkit for builders who need transparent checks for agent tool use, verification, and failure modes.

Grant ask: \$25,000 • Maintainer: Carly17



Why now

AI agents are moving from chat into real actions: code, files, infrastructure, APIs, and protocols.

The new risk

An agent can take action, skip verification, and still report success.

Small builders

Independent teams cannot use expensive closed evaluation systems.

Open gap


Safety rules need to be inspectable, adaptable, and locally runnable.

Failure pattern

action -> no verification -> false success -> user trusts the result

What the toolkit does

A local CLI evaluates agent traces and returns evidence-based safety findings.



False success

success without evidence

Missing verification

side effects without readback

Secret exposure

keys, tokens, env, wallets

Unsafe network writes

POST/PUT/DELETE without allowlist

Runnable proof

The repository includes code, tests, examples, CI, and a working demo.

Terminal

```
$ oask run unsafe_false_success.json
```

Agent: deployed and everything is working.

FAIL

Risk score: 30/100

false-success-without-evidence

5

tests passing

CI

GitHub Actions

MIT

open-source

28

repo files

Why this is a public good

Rules, traces, scoring, docs, and examples are open so builders can inspect and improve them.

Who benefits

Solo builders, open-source maintainers, small AI teams, web3 builders, and developers in emerging markets.

What stays open

Rules, example traces, scoring, documentation, CLI, tests, and integration examples.

What gets worse if closed

Small builders depend on hidden private benchmarks they cannot audit or adapt.

Core thesis

Agent safety should be something builders can run locally, not only something platforms sell privately.

Grant plan

A \$25,000 grant turns the prototype into a stronger public benchmark and toolkit.

1

Month 1 - Benchmark foundation

20+ curated traces, stable rule IDs, redaction guide, CI docs

2

Month 2 - Integrations

Transcript adapters, web3 checks, infrastructure checks, case studies

3

Month 3 - Public v0.1

50+ trace benchmark, downstream templates, benchmark report, release

Outcome: a usable open benchmark that helps builders verify agent actions before trusting them.