

Open Agent Safety Kit

Open-source safety tests for AI agents before real-world deployment.

Problem

AI agents increasingly run commands, edit files, deploy services, and interact with protocols. Small builders need practical ways to check whether agents verified their actions before claiming success.

Solution

A local CLI and open benchmark that evaluates agent traces for false success, missing verification, dangerous shell commands, secret exposure, unsafe network writes, hallucinated file claims, and web3 receipt gaps.

Why open

Closed evaluation makes independent builders depend on hidden rules. Open rules, traces, and scoring let builders inspect, adapt, and improve the safety layer together.

Current proof

- Runnable Python CLI: ``oask``
- Unit tests and GitHub Actions CI
- Safe and unsafe example traces
- MIT license and contribution/security docs
- Public repo: github.com/Carlys17/open-agent-safety-kit

Grant request

25,000 USD to expand the prototype into a public benchmark and toolkit: 50+ traces, transcript adapters, web3/infrastructure safety packs, CI templates, docs, and case studies.